

Lab 0: Environment Verification and First Measurement

Goal: Verify your environment works and run your first reproducible measurement.

Overview

This lab is started in class (Session 2) and completed as homework. By the end, you'll have:

- A working Ubuntu VM
- Basic tools installed (gcc, perf, strace)
- Your first performance measurement

Structure (Tiered Difficulty):

Part	Difficulty	Required?	Estimated Time
Part A	Basic	<input checked="" type="checkbox"/> Yes	30-45 min
Part B	Intermediate	<input checked="" type="checkbox"/> Yes	20-30 min
Part C	Advanced	 Optional	30-45 min

Deadline: Before Week 2 Lecture

Part A: Environment Setup (Basic — Required)

Step 1: Get Ubuntu VM Running

If you don't have an Ubuntu VM yet:

Recommended: VirtualBox (or VMware/UTM/Parallels)

- Download VirtualBox: <https://www.virtualbox.org/>
- Download an Ubuntu LTS ISO (recommended: **Ubuntu 22.04 LTS** or **Ubuntu 24.04 LTS**)
- Create VM: 2+ cores, 4GB+ RAM, 25GB+ disk
- Install Ubuntu, reboot
- See `README_VirtualBox_Setup.md` for detailed instructions

Alternative: Native Linux

- If you have Linux on bare metal, that works too

Kernel version requirement (do not skip)

Many later labs rely on modern kernel features (cgroup v2 details, eBPF + BTF). As a rule of thumb:

- **Minimum:** Linux kernel **5.10+**
- **Recommended:** Linux kernel **5.15+** (this is the default GA kernel series for Ubuntu 22.04 LTS)

Ubuntu guidance (so you do not accidentally pick an incompatible OS image):

- **Ubuntu 22.04 LTS** ships with a **5.15** kernel series by default (GA kernel).
- **Ubuntu 24.04 LTS** ships with a newer kernel series (also fine for this course).
- **Ubuntu 20.04 LTS** often runs kernel **5.4** by default (still 5.x, but may be missing features we rely on).

References:

- Ubuntu kernel lifecycle / GA vs HWE overview: <https://ubuntu.com/kernel/lifecycle>

You will also verify your actual kernel with `uname -r` below.

⚠ NOT Supported (as your main course environment):

- WSL2:** Many labs require kernel-level features that don't work reliably in WSL
- Docker containers:** You cannot access/modify the host kernel properly (perf/eBPF/cgroups are typically blocked)
- macOS:** Different kernel, no /proc, no perf

Note on Docker: *Installing Docker Engine inside your Ubuntu VM is required later (Weeks 6–7) for kind/Kubernetes labs.*

Step 2: Install Basic Tools

Open a terminal in Ubuntu and run:

```
# Update package list
sudo apt update

# Install build tools
sudo apt install -y build-essential

# Install perf
sudo apt install -y linux-tools-common linux-tools-$(uname -r)

# Install strace
sudo apt install -y strace

# Install git and curl
sudo apt install -y git curl wget
```

Step 2.5: Install "Stage 2" tools (required by Week 2)

Install these early so you do not get stuck in Week 2 labs:

```
sudo apt update
sudo apt install -y \
    python3 \
    valgrind
```

python3 is used by multiple labs for parsing logs and generating CSV summaries. valgrind (massif + ms_print) is required by Week 2A.

Step 3: Verify Installation

```
# Check GCC
gcc --version
```

```
# Check perf  
sudo perf stat ls  
  
# Check strace  
strace --version
```

All commands should work without "command not found" errors.

Step 4: Run Environment Check Script

```
# Run the course environment checker (provided in week1/)  
bash env_check.sh | tee lab0_env_check.txt
```

Review the output:

- [OK] for gcc/make/perf → Good!
- [WARN] for missing advanced tools (eBPF / Kubernetes / security) → OK for now, we will install them in later stages
- [ERROR] → Fix before continuing

Part A Checklist

- Ubuntu VM boots and is accessible
- gcc --version works
- sudo perf stat ls works
- strace --version works
- env_check.sh runs with no critical errors

Part A2: Course Toolchain (Staged Installs)

Lab 0 only verifies the minimum needed for Week 1. Later labs require additional packages and kernel features. We provide staged install blocks so you can install *only when needed*.

Stage 3 (install by Week 4): eBPF labs

These packages are used in Week 4+ labs that observe the scheduler with eBPF and BTF.

```
sudo apt update  
sudo apt install -y \  
    clang llvm \  
    libbpf-dev \  
    bptool \  
    bpftrace \  
    linux-headers-$(uname -r) \  
    stress-ng
```

Quick checks:

```
# BTF should exist on modern Ubuntu kernels (required for some BPF workflows)  
ls /sys/kernel/btf/vmlinux
```

```
# bpftrace should run (may require sudo)
sudo bpftrace -V
```

Stage 4 (install by Week 6): Kubernetes/kind labs

These labs require **Docker Engine inside your Ubuntu VM** (not "Docker as your OS environment"). We will provide exact install steps later to avoid distro/architecture pitfalls.

At minimum you will need:

- Docker Engine
- kind
- kubectl

Stage 5 (install by Week 8): storage observation tools

```
sudo apt update
sudo apt install -y sysstat

# iostat should be available
iostat -V
```

Stage 6 (install by Week 9): security labs

```
sudo apt update
sudo apt install -y \
    libcap-dev libcap2-bin \
    libseccomp-dev libseccomp2 \
    auditd
```

Environment feature checks (you will reuse these later)

```
# Kernel version (course minimum 5.10+, recommended 5.15+)
uname -r

# cgroup v2 (required for several labs)
mount | grep cgroup2 || true

# Dropping caches (used in storage/I/O labs; disruptive)
# sudo sysctl vm.drop_caches=3
```

Part B: First Measurement (Intermediate — Required)

Goal

Learn to use `perf stat` and understand what the numbers mean.

Step 1: Create a Simple Program

Create file `hello.c` :

```
#include <stdio.h>
#include <unistd.h>

int main() {
    printf("Hello from process %d\n", getpid());
    return 0;
}
```

Compile:

```
gcc -o hello hello.c
./hello
```

Step 2: Trace System Calls

```
strace ./hello
```

You'll see something like:

```
execve("./hello", ["../hello"], 0x7ffd...) = 0
...
write(1, "Hello from process 12345\n", 25) = 25
exit_group(0)                      = ?
```

Question: What system call does `printf` use under the hood?

Step 3: Measure with perf stat

```
sudo perf stat ./hello
```

Output looks like:

```
Performance counter stats for './hello':
```

0.42 msec	task-clock
1	context-switches
0	cpu-migrations
54	page-faults
912,345	cycles
456,789	instructions

Step 4: Understand the Numbers

Metric	What It Means
--------	---------------

task-clock	CPU time used by your program
context-switches	Times OS paused your program
page-faults	Memory pages allocated or loaded
cycles	CPU clock cycles consumed
instructions	Machine instructions executed

Step 5: Record and Explain

Run `perf stat 3 times` and record:

Run	task-clock (ms)	page-faults	cycles	instructions
1				
2				
3				

Questions to answer:

1. Are the numbers consistent across runs? Why or why not?
2. What is the Instructions Per Cycle (IPC = instructions / cycles)?
3. Why are there page faults for such a simple program?

Part B Checklist

- Compiled and ran `hello.c`
- Ran `strace ./hello` and saw system calls
- Ran `perf stat ./hello 3 times`
- Recorded results in a table
- Answered the explanation questions

Part C: Context Switch Experiment (Advanced — Optional)

Goal

Measure how context switches scale with workload and explain why.

Step 1: Create Test Program

Create file `ctx_switch_test.c`:

```
#include <stdio.h>
#include <stdlib.h>
#include <unistd.h>
#include <sys/wait.h>

int main(int argc, char *argv[]) {
    int n = 1000;
```

```

if (argc > 1) n = atoi(argv[1]);

int pipefd[2];
if (pipe(pipefd) < 0) {
    perror("pipe");
    return 1;
}

pid_t pid = fork();
if (pid < 0) {
    perror("fork");
    return 1;
}

if (pid == 0) {
    // Child: read from pipe n times
    close(pipefd[1]);
    char buf;
    for (int i = 0; i < n; i++) {
        if (read(pipefd[0], &buf, 1) < 0) {
            perror("read");
            exit(1);
        }
    }
    close(pipefd[0]);
    exit(0);
} else {
    // Parent: write to pipe n times
    close(pipefd[0]);
    char buf = 'x';
    for (int i = 0; i < n; i++) {
        if (write(pipefd[1], &buf, 1) < 0) {
            perror("write");
            exit(1);
        }
    }
    close(pipefd[1]);
    wait(NULL);
}

return 0;
}

```

Step 2: Compile

```
gcc -O2 -o ctx_switch_test ctx_switch_test.c
```

Step 3: Measure

```

sudo perf stat -e context-switches,cpu-migrations ./ctx_switch_test 100
sudo perf stat -e context-switches,cpu-migrations ./ctx_switch_test 1000
sudo perf stat -e context-switches,cpu-migrations ./ctx_switch_test 10000

```

Step 4: Record Results

n	context-switches	cpu-migrations	wall time (s)
100			
1000			
10000			

Step 5: Explain

Answer these questions:

1. **Scaling:** How does context-switch count scale with n? Is it linear?
2. **Mechanism:** Why does this program cause context switches?
 - o Hint: What happens when child calls `read()` but no data is available?
 - o Hint: What happens when parent calls `write()` but buffer might be full?
3. **Prediction:** What would happen if we used larger buffer sizes?

Part C Checklist

- Created and compiled `ctx_switch_test.c`
- Measured with n = 100, 1000, 10000
- Recorded results in table
- Explained the scaling behavior and mechanism

Submission

Create a file `lab0_report.md` with this structure:

```

# Lab 0 Report

**Name:**  

**Date:**  
  

## Part A: Environment Setup  
  

### Environment Check Output  

(Paste key lines from env_check.sh output)  
  

### Issues Encountered  

(Describe any problems and how you fixed them, or "None")  
  

## Part B: First Measurement

```

```
### strace Output  
(Paste the key system calls you observed)
```

****Question:**** What system call does printf use?

****Answer:****

```
### perf stat Results
```

Run	task-clock (ms)	page-faults	cycles	instructions
1				
2				
3				

```
### Explanation Questions
```

1. Are the numbers consistent? Why?

2. What is the IPC?

3. Why are there page faults?

```
## Part C: Context Switch Experiment (Optional)
```

(Include your table and explanations if completed)

Submit: lab0_report.md before Week 2 Lecture

Troubleshooting

"perf: command not found"

```
sudo apt install -y linux-tools-common linux-tools-$(uname -r)
```

"Permission denied" with perf

```
# Option 1: Use sudo  
sudo perf stat ./program  
  
# Option 2: Lower security level (for your own VM only)  
sudo sysctl kernel.perf_event_paranoid=1
```

perf shows "not supported"

Some hardware counters may not be available in VMs. That's OK — basic counters (cycles, instructions, task-clock) should work.

VM is very slow

- Allocate more RAM (4GB minimum, 8GB recommended)
- Allocate more CPU cores (2 minimum)
- Enable VT-x/AMD-V in BIOS settings

Can't install linux-tools for your kernel version

```
# Check your kernel version  
uname -r  
  
# Install generic tools  
sudo apt install -y linux-tools-generic
```

Getting Help

- **During lab workshop:** Ask the instructor or TA
- **Outside of class:** Office hours, course forum
- **Environment issues:** Don't struggle alone — ask early!

Environment problems get harder to fix under deadline pressure. Get help now.

Grading Rubric

Criterion	Points
Part A (30 points)	
Environment check passes	20
Issues documented clearly	10
Part B (60 points)	
strace output included	10
perf stat results (3 runs)	20
Explanations are correct and thoughtful	30
Part C (10 bonus points)	
Experiment completed	5
Mechanism explained correctly	5

Total: 90 points (+ 10 bonus)