

# **Advanced Operating Systems**

**For the Cloud-Native & AI Era**

**Week 1: What Problems Does This Course Solve?**

# This Week's Plan

## Session 1 (80 min): Course Introduction

- Who are we?
- What is this course about?
- How will you be evaluated?
- What do you need to set up?

## Session 2 (80 min): Environment Setup

- Get your tools working
- Run your first measurement

# Who Am I?

**Dong Dai**

- Associate Professor, Computer & Information Sciences, University of Delaware
- PhD from University of Science and Technology of China (USTC)

**Research Areas:**

- High-Performance Computing (HPC) systems
- Parallel file systems and I/O optimization
- AI for Systems (AI4HPC)

# Data Intelligence Research Lab (DIRLab)

## What we do:

- Optimize data-intensive infrastructure: parallel file systems, metadata management, job scheduling
- Apply machine learning to understand and predict system behavior
- Build intelligent tools for HPC storage optimization

## Recent work (2025):

- [SC'25] STELLAR: Using LLM autonomous reasoning to tune parallel file systems
- [SC'25] Improving SpGEMM performance through matrix-reordering
- [HPDC'25] TSUE: Two-stage data update for erasure-coded cluster file systems

# Teaching Assistant

**Minqiu Sun**

- Ph.D. Student, DIRLab Member
- Research: LLM Post-training and Agent System Design
- [PDSW'25] LLMTailor: Layer-wise tailoring for efficient LLM checkpointing

**Office Hours:** TBA

**Responsibilities:** Lab support, grading, office hours

# Who Are You?

**Round table introduction:**

- Your name
- Your program (PhD / MS / Undergrad)
- Your background (what OS/systems courses have you taken?)
- One system you want to understand better

*(Take notes — you might find project partners here!)*

# Course Information

<b>Class Time</b>	12:45 – 2:05pm
<b>Location</b>	Memorial Hall, Room 113
<b>Lecturer</b>	Dong Dai ( <a href="https://daidong.github.io">daidong.github.io</a> )
<b>Office</b>	Fintech 416B
<b>TA</b>	TBA
<b>Office Hours</b>	TBA

# Course Policies: Attendance

This is an in-person course.

-  All students are required to attend every class
-  Sick leave requires a doctor's note
-  **20% of your grade** is participation
-  In-class quizzes and activities

*If you're not in class, you can't participate!*

# Course Policies: Assignments

## Assignments:

- Most are **hands-on labs** (coding + measurement + explanation)
- May include some written assignments

## Deadlines & Late Submission:

-  All deadlines are **final**
-  In case of unforeseen events: **3-day safety margin** (use wisely!)
-  **NO late submissions** will be accepted after the margin

*Plan ahead. Start early.*

# Course Policies: Projects & Exams

## Programming Projects:

- Final project is a system case study (40% of grade)
- Can work individually or in pairs
- Milestones throughout the semester

## Exams:

 **NO EXAMS!**

The text "NO EXAMS!" is centered, flanked by two small party hats with streamers. The hats are red and green with yellow streamers.

Your work (labs + project) speaks for itself.

# Academic Integrity

## Individual Work:

- All assignments are to be completed **individually** (unless specified)
- **✗** No code from other students
- **✓** Code from generative AI (ChatGPT, Copilot, etc.) is **welcome**

## Collaboration & AI Use:

- **✓** You are encouraged to **discuss** with classmates
- **✓** You are encouraged to use **ChatGPT** to facilitate learning
- **⚠** But you must **understand** what you submit

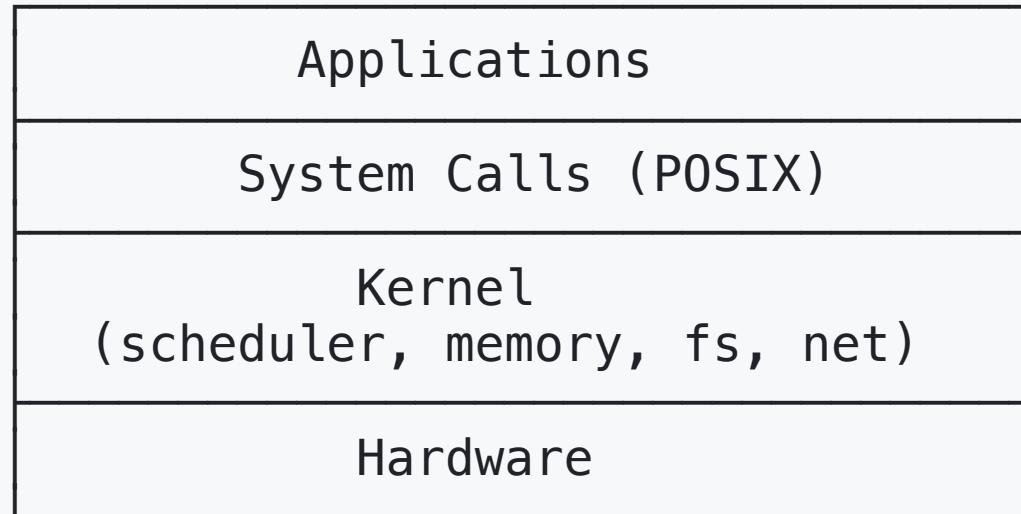
*If you can't explain your code, you didn't write it.*

# The Million Dollar Question

What is an "Operating System" in 2025?

- Is it just the Linux kernel?
- Where does the OS end and the "platform" begin?
- What should a graduate student learn about OS today?

# Traditional OS Course



Focus: internals of a single kernel on a single machine.

**This remains essential.** We will cover these core concepts.

# OS Core Concepts We Will Cover

Concept	What You'll Learn
Process & Thread	Creation, lifecycle, context switch
Scheduling	Policies, fairness, latency
Virtual Memory	Paging, TLB, page cache
Concurrency	Locks, synchronization, races
I/O & Storage	Filesystems, buffering, persistence
Security	Isolation, permissions, sandboxing

These are **not optional**. You need them to understand what's happening.

# What's Different in Practice Today?

The core concepts remain, but the **context** has changed:

- **Containers**: process isolation using namespace + cgroup
- **Kubernetes**: cluster-level resource management
- **eBPF**: kernel observability without modifying kernel code
- **New workloads**: LLM inference, AI agents

We'll use these as **case studies** to apply OS concepts.

# This Course: A Perspective

Control Plane  
(K8s scheduler, controllers, orchestration)

Runtimes  
(container runtime, eBPF programs)

Kernel  
(scheduler, memory, fs, net, isolation)

Hardware

This is **one way** to view modern systems. Not the only way.

# Course Thesis

OS core concepts (processes, memory, scheduling, I/O, security) are the foundation.  
We learn them by applying them to real systems and measuring real behavior.

This course trains you to:

1. **Understand** — Master OS fundamentals: processes, memory, scheduling, I/O, isolation
2. **Measure** — Design and run reproducible experiments with real metrics
3. **Explain** — Trace observed behavior back to specific OS mechanisms

# Why This Approach?

The skills you need to debug a "p99 spike" or "container OOM" are built on OS internals:

Symptom	OS Concept	Tool
p99 latency spike	Scheduling, page fault	perf, eBPF
Container OOM	Memory management, cgroups	cgroup stats
CPU throttling	CFS bandwidth control	/proc, /sys
Network delay	TCP/IP stack, context switch	tcpdump, bpftrace

Understanding **why** requires knowing the OS.

# Learning Objectives

By the end of this course, you will be able to:

1. **Explain** core OS concepts: process lifecycle, scheduling policies, virtual memory, file systems, and concurrency primitives
2. **Use measurement tools** (perf, eBPF, strace) to investigate performance issues with reproducible experiments
3. **Construct isolation boundaries** (namespace, cgroup, seccomp) and explain the underlying OS mechanisms
4. **Analyze resource management** in container and cluster environments (reproduce OOM, throttling, eviction)
5. **Profile modern workloads** (LLM inference, agent tools) and trace bottlenecks to OS-level causes

# Two Sessions Every Week

Session	What Happens	Graded?
<b>Session 1: Lecture</b>	I teach concepts + case studies	In-class quiz → Participation (20%)
<b>Session 2: Lab Workshop</b>	You start the lab, I help	Lab submission → Labs (40%)

Both sessions are **required**. Missing lecture = missing participation points.

# How It All Fits Together

## SESSION 1: LECTURE (in class, 80 min)

- I teach OS concepts + modern context + case study
- Preview what you'll do in the lab
-  In-class quiz/activity → Participation (20%)

↓

## SESSION 2: LAB WORKSHOP (in class, 80 min)

- You START the lab with me and TA available
- Get help, ask questions, understand requirements
-  You will NOT finish the lab in class!

↓

## AFTER CLASS: Complete Lab Assignment

- Continue working on the lab outside of class
- Deadline: before next week's lecture
-  Submit script + data + explanation → Labs (40%)

# Lecture: What We'll Cover

Each lecture will include:

- 1. OS concepts** (30-40 min): The fundamentals you need to know
- 2. Modern context** (20-30 min): How these concepts apply today
- 3. Case study** (15-20 min): A real scenario demonstrating the concepts
- 4. Lab preview** (10 min): What you'll be measuring and why

I'll provide reading materials, but the lecture is where we build shared understanding.

# Lab Workshop: What Happens in Class

The lab workshop is for **getting started**, not finishing:

1. **Warm-up** (10 min): Review objectives and metrics
2. **Guided Build** (30 min): I show the skeleton and common pitfalls
3. **Work Time** (30 min): You start working, ask questions
4. **Debrief** (10 min): Share progress, identify common issues

**After class:** You continue working and submit before next lecture.

# Lab Assignment: What You Submit

After each lab workshop, you complete the lab and submit:

- 1. Reproducible scripts (Makefile or bash)**
- 2. Data files (CSV, logs, measurements)**
- 3. Explanation document (1-2 pages)**

**Deadline:** Before next week's lecture (exact time on each lab handout)

**~8 labs total**, each worth ~5% of your grade.

# What Makes a Good Lab Report?

A good lab report demonstrates:

- 1. Reproducibility:** I can run your script on a clean VM and get similar results
- 2. Clear metrics:** Numbers have units, sampling windows, and defined baselines
- 3. Comparison:** At least one baseline (before vs. after, config A vs. B)
- 4. Causal explanation:** Link observed behavior to a specific OS mechanism

## Lab Report Anti-Patterns

- ✗ "It was slow" — What does "slow" mean? Compared to what?
- ✗ "We ran the experiment" — Can I reproduce it? Where's the script?
- ✗ "The latency increased" — By how much? Why? What OS mechanism?
- ✓ "p99 latency increased from 2ms to 15ms when memory limit was reduced to 512MB, because page reclaim (kswapd) was triggered every 100ms"

# A Note on Student Diversity

This class has PhD students, Master's students, and undergraduates.

## My approach:

- Lectures will cover concepts clearly — no one gets left behind
- Labs will have tiered difficulty (basic requirements + advanced challenges)
- Final project scope can be adjusted based on your level
- Ask questions! There are no dumb questions.

# Final Project: System Case Study

Each student (or pair) picks a "**system case**" — a real system exhibiting a behavior worth investigating.

Examples:

- PostgreSQL p99 spike under memory pressure
- Redis latency when competing with noisy neighbor
- nginx container OOM under bursty load
- llama.cpp inference timeout under concurrent requests

# Final Project Milestones

Week	Milestone	Deliverable
4	Proposal	1-page problem statement + minimal repro
6	Checkpoint	Tracing evidence (eBPF or perf)
10	Freeze	Clear metric + intervention plan
12	Defense	Presentation + reproduction by peers

Your project is **40% of grade**. Start thinking now.

# Grading Breakdown

Component	Weight	What It Includes
Final Project	40%	System case study (4 milestones)
Labs	40%	~8 lab assignments (~5% each)
Participation	20%	In-class quizzes, activities, discussion

**No closed-book exams.** Your work speaks for itself.

# What Counts as Participation?

Earned through **in-class activities during lectures**:

- In-class quizzes (unannounced, short)
- In-class activities and discussions
- Asking and answering questions
- Helping classmates

**You must attend lecture to earn participation credit.**

# 12-Week Roadmap (Preview)

Week	Theme	OS Concepts
1	Introduction	Setup + methodology
2	Performance	Memory hierarchy, cache, measurement
3	Concurrency	Threads, scheduling policies
4	eBPF	Tracepoints, kernel observability
5	Containers	Namespaces, cgroups
6	Kubernetes	Resource limits, QoS, eviction

# 12-Week Roadmap (continued)

Week	Theme	OS Concepts
7	Network(*)	TCP/IP stack, socket buffers
8	Storage	Filesystems, page cache, fsync
9	Security(*)	Capabilities, seccomp, sandboxing
10	LLM Workloads	Memory pressure, inference profiling
11	Agent Runtime	Tool isolation, least privilege
12	Final Defense	Presentations + peer reproduction

## Part 2: Environment Setup

We'll spend the rest of Session 1 walking through environment setup.

**Why this matters:**

- Labs require **kernel-level access** (perf, eBPF, cgroups)
- Different environments = different results = debugging nightmares
- We all need to be on the **same page**

# The Two-Environment Model

You'll work with **two environments**:

Environment	Purpose	What Runs Here
Your laptop	Writing code, reports, browsing	Editor, browser, PDF viewer
Ubuntu VM	Running labs, measurements	perf, eBPF, containers, experiments

**Why?** Your laptop (macOS/Windows) can't run kernel-level tools reliably.

## ⚠️ Important: What NOT to Use

Environment	Problem	Result
WSL2	Kernel features limited	eBPF/perf may fail silently
Docker	BPF capabilities restricted	Can't access PMU, limited /proc
macOS	Different kernel	No /proc, no perf, no cgroups

The only supported environment: Ubuntu VM (VirtualBox) or native Ubuntu.

# Step 1: Install VirtualBox

Download: [virtualbox.org](https://www.virtualbox.org)

For Intel/AMD (most Windows/Linux):

- Download the standard installer
- Follow the wizard

For Apple Silicon (M1/M2/M3 Mac):

- VirtualBox support is limited
- Consider **UTM** as an alternative
- Ask me or TA if you need help

## Step 2: Download Ubuntu ISO

**Recommended:** Ubuntu 22.04 LTS or Ubuntu 24.04 LTS

**Course requirement (hard):** Linux kernel 5.15+

- Check your kernel with: `uname -r`
- Ubuntu 22.04 LTS ships with a 5.15 kernel series by default

**Download links:**

- Intel/AMD: [ubuntu.com/download/desktop](https://ubuntu.com/download/desktop)
- ARM (Apple Silicon): [ubuntu.com/download/server/arm](https://ubuntu.com/download/server/arm)

**File size:** ~4-5 GB — download before class if possible!

## Step 3: Create the VM

In VirtualBox, click "New":

Setting	Recommended Value
Name	Ubuntu-OS-Course
Type	Linux
Version	Ubuntu (64-bit)
RAM	4 GB (8 GB if possible)
CPU	2 cores
Disk	25 GB or more

Then attach the Ubuntu ISO and install.

## Step 4: Install Ubuntu

1. Boot the VM with ISO attached
2. Follow the installer (defaults are fine)
3. Create a user account (remember your password!)
4. Wait for installation to complete
5. Remove ISO and reboot

**Time estimate:** 15-30 minutes

## Step 5: Install Development Tools

Open a terminal in Ubuntu and run:

```
# Update package list  
sudo apt update  
  
# Install basics  
sudo apt install -y build-essential git curl wget  
  
# Install perf  
sudo apt install -y linux-tools-common linux-tools-$(uname -r)  
  
# Install strace (system call tracing)  
sudo apt install -y strace
```

## Step 6: Verify Basic Tools

```
# Check GCC  
gcc --version  
  
# Check Make  
make --version  
  
# Check Git  
git --version  
  
# Check perf (needs sudo)  
sudo perf stat ls
```

- ✓ All should print version numbers or run without errors.

## Step 7: Install eBPF Tools (For Week 4+)

```
# eBPF tracing tools  
sudo apt install -y bpftrace bpfcc-tools  
  
# eBPF development libraries  
sudo apt install -y clang llvm libbpf-dev libelf-dev  
  
# Kernel headers (needed for eBPF compilation)  
sudo apt install -y linux-headers-$(uname -r)
```

**Note:** If these fail, it's OK — we don't need them until Week 4.

## Step 8: Install Optional Tools

```
# Stress testing tools (for generating workloads)
sudo apt install -y stress-ng fio

# Python for data analysis
sudo apt install -y python3 python3-pip
pip3 install matplotlib pandas numpy

# htop (better process viewer)
sudo apt install -y htop
```

# Tool Summary: What You'll Use

Tool	Purpose	When
gcc	Compile C code	Every lab
perf	Performance measurement	Week 2+
strace	Trace system calls	Week 2+
bpftrace	eBPF tracing	Week 4+
stress-ng	Generate CPU/memory load	Week 5+

## Step 9: Run the Course Environment Check Script

From the course repo (see Lab 0), run:

```
bash env_check.sh | tee lab0_env_check.txt
```

**Goal:** no **critical errors** (including kernel < 5.15).

# Self-Test: Is Your Environment Ready?

Run these commands. All should work:

```
# 1. Compile a C program
echo 'int main(){return 0;}' > test.c && gcc -o test test.c && ./test && echo "OK"

# 2. Run perf
sudo perf stat ls >/dev/null

# 3. Explore /proc
cat /proc/cpuinfo | head -5

# 4. Run strace
strace -c ls 2>&1 | head -5
```

# Self-Test Checklist

Check	Command	Status
GCC works	<code>gcc --version</code>	<input type="checkbox"/>
Make works	<code>make --version</code>	<input type="checkbox"/>
Perf works	<code>sudo perf stat ls</code>	<input type="checkbox"/>
/proc accessible	<code>cat /proc/cpuinfo</code>	<input type="checkbox"/>
strace works	<code>strace -c ls</code>	<input type="checkbox"/>

All checked = Ready for Lab 0!

# Your First C Program

Create `hello.c`:

```
#include <stdio.h>
#include <unistd.h>

int main() {
    printf("Hello from process %d\n", getpid());
    return 0;
}
```

Compile and run:

```
gcc -o hello hello.c
./hello
```

# Your First System Call Trace

```
strace ./hello
```

You'll see something like:

```
execve("./hello", ["./hello"], ...) = 0
...
write(1, "Hello from process 12345\n", 25) = 25
exit_group(0)                      = ?
```

This is what OS does behind the scenes!

# Your First Measurement: perf stat

```
sudo perf stat ./hello
```

Output:

```
Performance counter stats for './hello':  
    0.42 msec task-clock  
          1      context-switches  
          0      cpu-migrations  
         54      page-faults  
  912,345      cycles  
 456,789      instructions
```

You just measured your program's OS-level behavior!

# Understanding perf Output

Metric	What It Means
<b>task-clock</b>	CPU time used by your program
<b>context-switches</b>	Times OS paused your program
<b>page-faults</b>	Memory pages allocated or loaded (minor + major faults)
<b>cycles</b>	CPU clock cycles consumed
<b>instructions</b>	Machine instructions executed

This is the foundation of performance analysis!

# Lab 0: Submission (Canvas)

Submit ONE ZIP file on Canvas.

- **ZIP name (recommended):** `lab0_<lastname>_<firstname>.zip`
- **Due:** Before Week 2 lecture (see Canvas for the exact timestamp)

**Minimum requirement:** the ZIP must contain `lab0_report.md`.

# Lab 0: What Goes Into the ZIP?

File	Required?	What It Is For
lab0_report.md	Yes	Your answers + results (the graded report)
lab0_env_check.txt	Recommended	Raw output of <code>bash env_check.sh   tee lab0_env_check.txt</code>
hello.c	Recommended	Your “first program” source code (trace/measure target)
perf_stat_hello.txt	Optional	Raw <code>perf stat ./hello</code> output (supports reproducibility)
strace_hello.txt	Optional	Raw <code>strace ./hello</code> output (supports reproducibility)

# Lab 0: Grading Rubric (90 pts + 10 bonus)

Criterion	Points
<b>Part A (30 pts)</b>	
Environment check passes	20
Issues documented clearly	10
<b>Part B (60 pts)</b>	
strace output included	10
perf stat results (3 runs)	20
Explanations are correct and thoughtful	30
<b>Part C (10 bonus pts)</b>	
Experiment completed	5

# Troubleshooting: Common Issues

## "command not found"

→ Install the missing package: `sudo apt install <package>`

## "Permission denied" with perf

→ Use `sudo`: `sudo perf stat ...`

## perf shows "not supported"

→ You might be in a VM with limited PMU access — that's OK, some counters work

## "No kernel headers"

→ `sudo apt install linux-headers-$(uname -r)`

# Troubleshooting: VirtualBox Issues

## VM is very slow

→ Enable VT-x/AMD-V in BIOS, allocate more RAM

## No network in VM

→ Check VirtualBox network settings (NAT or Bridged)

## Screen resolution stuck

→ Install Guest Additions: `sudo apt install virtualbox-guest-utils`

## Can't copy/paste between host and VM

→ Install Guest Additions and enable clipboard sharing

# Getting Help

## During lab session:

- Ask me or TA — we're here to help!

## Outside of class:

- Check course materials first
- Post on course forum (if available)
- Office hours

**Don't struggle alone.** Environment issues are common and fixable.

## Week 2 Preview: What's Coming

**Topic:** Performance Methodology — Measurement to Mechanism

We will cover:

- **OS concepts:** Memory hierarchy, cache, TLB
- **Measurement:** What makes a valid experiment?
- **Tools:** perf, time, valgrind
- **Lab:** Analyze a sorting algorithm's performance

# Summary: What We Covered Today

1. **Course structure:** Lecture + Lab Workshop + Lab Assignment
2. **Grading:** 40% Project + 40% Labs + 20% Participation
3. **Environment:** Ubuntu VM with perf, strace, eBPF tools
4. **First measurement:** `perf stat` and `strace`

**Session 2:** Let's actually set up your VMs!

# Questions?

- About the course?
- About the project?
- About the environment?

**Let's get your VMs set up!**