

Chapter 7 Scheduling

Dong Dai

CIS/UD

dai@udel.edu

Scheduling

- CPU Virtualization
 - low-level mechanisms - context switching
 - high-level policies - **scheduling policies**
- Appears in many fields - e.g., factory assembly lines for efficiency
- How to develop scheduling policies? How to develop a basic framework for studying scheduling policies
 1. assumptions of workload
 2. metrics of performance
 3. approaches used

1. Workload Assumptions


Workload - {processes} running in the system

- critical to building policies - the more you know, the more fine-tuned your policy can be
 - unrealistic, initially, but will be relaxed as we go
1. each job runs for the same amount of time
 2. all jobs arrive at the same time
 3. once started, each job runs to completion
 4. all jobs only use CPU (i.e., they perform no I/O)
 5. the run-time of each job is known

Two kinds: batch (CPU-bound) vs. interactive (I/O-bound)


2. Scheduling Metrics

- To compare different scheduling policies, we need **metrics** to measure behavior and hence performance
- Turnaround time - the time at which the job completes minus the time at which the job arrived in the system
 - a **performance** metric
- Fairness metric
 - Jain's Fairness Index
- Performance vs. fairness **tradeoff**
 - a scheduler may optimize performance but at the cost of preventing a few jobs from running, thus decreasing fairness



First In First Out (FIFO)

- First come, first served (FCFS)



Workload assumptions:


1. Each job runs for the same amount of time
2. All jobs arrive at the same time
3. Once started, each job runs to completion
4. All jobs only use CPU (i.e., they perform no I/O)
5. The run-time of each job is known

- Each job runs 10 sec.
- Average turnaround time = ?
 $= (10+20+30)/3 = 20$

turnaround time = completion time - arrival time

First In First Out (FIFO)

- What kind of workload could you construct to make FIFO perform **poorly**?
- Relax assumption #1: jobs could have different (still known) run times




- Average turnaround time = ?
 $= (100+110+120)/3 = 110$
- Convoy effect (supermarket checkout line)
- How can you improve it? (new scheduling policy?)

Workload assumptions

- 1. Each job runs for the same amount of time
- 2. All jobs arrive at the same time
- 3. Once started, each job runs to completion
- 4. All jobs only use CPU (i.e., they perform no I/O)
- 5. The run-time of each job is known

Shortest Job First (SJF)

- Average turnaround time = ?
 $= (10+20+120)/3 = 50$



Workload assumptions

- 1. Each job runs for the same amount of time
- 2. All jobs arrive at the same time
- 3. Once started, each job runs to completion
- 4. All jobs only use CPU (i.e., they perform no I/O)
- 5. The run-time of each job is known

- Assuming all jobs arrive at the same time, SJF is optimal


- Relax assumption #2: arrival time

average turnaround time

$$= (100+(110-10)+(120-10))/3$$

$$= 103.33$$

- How can we do better?




Shortest Time-to-Completion First

- SJF with A arrives at 0, and B, C at 10

Workload assumptions

1. Each job runs for the same amount of time
2. All jobs arrive at the same time
3. Once started, each job runs to completion
4. All jobs only use CPU (i.e., they perform no I/O)
5. The run-time of each job is known




- How would you do better?

Relax assumption #3


→ **preemption**

→ **STCF**



or **preemptive SJF** (also optimal with assumption that jobs arrive at any time)

- What if we can not **Relax assumption #3?**



Delay Scheduling*

SchedInspector: A Batch Job Scheduling Inspector Using Reinforcement Learning

Di Zhang

Computer Science Department,
University of North Carolina at
Charlotte
Charlotte, NC, USA
dzhang16@uncc.edu

Dong Dai

Computer Science Department,
University of North Carolina at
Charlotte
Charlotte, NC, USA
ddai@uncc.edu



Bing Xie

Oak Ridge Leadership Computing
Facility, Oak Ridge National
Laboratory
Oak Ridge, TN, USA
xieb@ornl.gov


ACM HPDC 2022

The 31st International Symposium on High-Performance Parallel and Distributed Computing
Minneapolis, Minnesota, United States, June 27 - July 1, 2022

Delay Scheduling*



(a) Sufficient resources case



(b) Insufficient resources case

Figure 1: Scheduling jobs with/without SchedInspector. *x-axis shows the timeline in minutes; y-axis shows the compute nodes ($n_1 \rightarrow 5$); each block represents a job that takes amount of nodes and time; J_p is the preliminary job running before the scheduling starts.*

Delay Scheduling*








Figure 3: Architecture of SchedInspector.

Delay Scheduling*

(a) *wait*



(b) *mbsld*



New Metric


Workload assumptions

1. Each job runs for the same amount of time
2. All jobs arrive at the same time
3. Once started, each job runs to completion
4. All jobs only use CPU (i.e., they perform no I/O)
5. The run-time of each job is known


- With known job lengths, **CPU-only** jobs, and metric of turnaround time, STCF is great for **batch** systems (processes that run without user interaction)
- **Time sharing** system with **interactive** applications - need new metric
- **Response time** = time from when the job arrives in a system to the first time it is scheduled

Response Time

- Response time = $(0 + 0 + 10)/3 = 3.33$





- Is STCF (or SJF) good for interactive applications?
 - not at all (with just preemption)!
- What can we do?



Round Robin (RR)

- Instead of running jobs to completion, RR runs a job for a (small) **time slice** (scheduling quantum) and switches to the next job in the ready queue; repeatedly does so until jobs are finished
- Time slicing** - **length of time slice** = multiple of timer-interrupt period
 - SJF avg. response time = 5
 - RR avg. response time = 1
- Length of time slice is critical**
 - under response time, the shorter the better
 - too short, overhead of context switching dominates



Round Robin

- Amortization (in general)
 - used in systems when there is a fixed cost to some operation (e.g., context switching)
 - by incurring that cost **less often**, the total cost to the system is reduced
- For example, if time slice == 10 ms, and the context-switch cost is 1 ms, roughly 10% of time is spent context switching and is thus wasted
- To amortize this cost, we can increase the time slice, to, e.g., 100 ms. In this case, less than 1% of time is spent context switching, and thus the cost of time-slicing has been amortized

Round Robin

- Tradeoff?
- Length of time slice on response time ???
 - the shorter the time slice, the **better** the response
 - too short, cost of context switching (**overhead**) dominates
 - length of time slice is a **tradeoff** - long enough to amortize the cost of context switching without making it so long that the system is no longer responsive
- In addition, any policy (such as RR) that is **fair** (evenly divides CPU among active processes on a small time scale) will perform poorly on metrics such as turnaround time

So Far...


- Two types of schedulers
 - SJF and STCF: optimize **turnaround** time, but is bad for response time
 - RR: optimize **response** time, but is bad for turnaround time

Workload assumptions

1. Each job runs for the same amount of time
2. All jobs arrive at the same time
3. Once started, each job runs to completion
4. All jobs only use CPU (i.e., they perform no I/O)
5. The run-time of each job is known

Incorporating I/O

- Process is blocked waiting for I/O completion
- When I/O completes, an interrupt is raised, and OS runs and moves the process that issued the I/O from blocked state back to ready (or even running) state
- Poor use of CPU
- Higher CPU utilization?
 - overlap CPU & I/O



- Interactive jobs get run frequently; while they are performing I/O, other CPU-intensive jobs run

No More Oracle

1. Each job runs for the same amount of time
2. All jobs arrive at the same time
3. Once started, each job runs to completion
4. All jobs only use CPU (they perform no I/O)
5. The run-time of each job is known
 - how can we build a scheduler that behaves like SJF/STCF **without** *a priori* knowledge of runtime?
Workload assumptions
 1. Each job runs for the same amount of time
 2. All jobs arrive at the same time
 3. Once started, each job runs to completion
 4. All jobs only use CPU (i.e., they perform no I/O)
 5. The run time of each job is known

next chapter on multi-level feedback queue



Can you answer these?

- True/False
 - A timer tick is identical to a time slice
 - The convoy effect occurs when longer jobs must wait for shorter jobs
 - A RR scheduler may preempt a previously running job
- Essay Questions
 - Assume the OS schedules a workload containing three jobs with the following characteristics:

Job	Arrival Time	CPU Time
A	0	10
B	5	8
C	12	2

- Which scheduler minimizes the completion time of the entire workload? FIFO, RR, SJF, STCF?
- Given a **STCF** scheduler, what is the response time for Job B?
- Given a **STCF** scheduler, what is the response time for Job C?



Can you answer these?

- True/False
 - A timer tick is identical to a time slice **[False]**
 - The convoy effect occurs when longer jobs must wait for shorter jobs **[False]**
 - A RR scheduler may preempt a previously running job **[True]**
- Essay Questions
 - Assume the OS schedules a workload containing three jobs with the following characteristics:

Job	Arrival Time	CPU Time
A	0	10
B	5	8
C	12	2
 - Which scheduler minimizes the completion time of the entire workload? FIFO, RR, SJF, STCF?
 - None of the above. All schedulers give the same results. The completion time for the entire workload (i.e., the last job) will be the same independent of the completion time of the individual jobs
 - Given a **STCF** scheduler, what is the response time for Job B?
 - B is scheduled at time 10 (at time 5 when B arrives, it has a burst of 8 > A's remaining burst of 5); $10-5 = 5$
 - Given a **STCF** scheduler, what is the response time for Job C?
 - When C arrives at time 12, it has a CPU burst of 2 < B's remaining CPU burst of $8-2 = 6$; therefore, C is scheduled immediately when it arrives.